



KOGOD SCHOOL *of* BUSINESS
AMERICAN UNIVERSITY • WASHINGTON, DC

ITEC 621 Predictive Analytics Project

Project Name:
Predicting Consumer Behavior: Insights from E-Commerce Data

Class Section: 002

Team Number: 7

Team Members:

Marco Capoccia, Viraj Pande, Dheeraj Shetty, Victoria Smith

Last updated: May 1, 2025

Deliverable Number: 4

1. Business Case

As online retail expands, companies face ongoing challenges of attracting website visitors and ensuring they complete a purchase. One of the largest challenges is determining whether the investment in acquiring and retaining customers is profitable, especially given that average e-commerce conversion rates are only around 2.63% (Smith, 2024). That being said, it is critical to understand key behavioral drivers in online shopping for companies to grow revenue and ROI. Although there are platforms available to provide behavioral metrics, it is challenging for most businesses to turn this data into actionable insights to help support strategic business decision-making.

This challenge is also magnified by the rising cost of acquiring new customers. According to research, customer acquisition costs (CAC) have increased by over 220% in the past eight years, starting from an average of \$9 per customer (2013) to around \$29 (2022) (O'Rourke, 2022). In a competitive environment like this, even minor improvements in conversion rates can result in substantial revenue gains and more efficient marketing spend.

This project aims to fill the gap by analyzing data from over 12,000 buying sessions, each showing whether a shopper made a purchase. By examining variables like time spent on the product page, type of visit, and time of visit (weekend vs. weekday), we aim to discover which variables are most useful in predicting purchase behavior. We will use statistical models to predict the likelihood of a purchase and identify which behavior traits matter most. The results of this analysis can help e-commerce retailers better understand user shopping patterns and increase conversion rates and profits.

2. Business and Analytics Questions

2.1. Business Question: How can an online retailer increase sales by identifying key behavioral factors that drive shoppers to complete a purchase?

2.2. Analytics Question: What is the impact of session behavior (e.g., time spent on product pages, bounce rates, exit rates) and visitor traits (e.g., returning customers, weekend visits) on the likelihood of a purchase? Additionally, how does the month of visit influence purchase probability?

Analytics Goal: Our goal is to identify factors that predict online purchases and provide actionable insights for e-commerce optimization. Using classification models like Logistic Regression and Random Forests, we aim to balance predictive accuracy with interpretability for stakeholders.

3. Data Set Description

Our analysis uses the "Online Shoppers Intention" dataset, comprising 12,330 buying sessions from an e-commerce platform, each indicating whether a purchase occurred. The dataset is sourced from the UCI Machine Learning Repository. The main variables of interest are:

Outcome Variable: Revenue (binary: True/False) – Indicates whether a purchase was made. The dataset is unbalanced, with 10,422 non-purchase sessions (False) and 1,908 purchase sessions (True).

Predictors:

ProductRelated_Duration (numeric, seconds) – Time spent on product pages.

BounceRates (numeric, proportion) – Proportion of single-page visits.

ExitRates (numeric, proportion) – Proportion of sessions ending on a page.

VisitorType (categorical: Returning/New/Other) – Visitor status.

Weekend (binary: True/False) – Whether the visit occurred on a weekend.

Month (categorical: e.g., Feb, Mar) – Month of the visit.

These variables were chosen for their relevance to online shopping behavior and potential to influence purchase decisions.

4. Descriptive Analytics

4.1. Descriptive Statistics of Focal Variables

To understand the baseline characteristics of our predictors, we computed means and standard deviations for numeric variables and frequency distributions for categorical variables. For instance, shoppers who made a purchase spent significantly more time on product pages (mean = 1,876 seconds) compared to non-buyers (mean = 1,070 seconds), suggesting engagement as a potential driver of purchases. Similarly, bounce rates and exit rates were lower for buyers (means = 0.005 and 0.020, respectively) than non-buyers (means = 0.025 and 0.047), indicating that less site abandonment correlates with purchases.

4.2. Distribution of Outcome and Other Variables

The outcome variable, Revenue, is binary (True/False), representing whether a purchase occurred. A bar plot shows an unbalanced distribution: 84.5% of sessions resulted in no purchase (10,422), while 15.5% led to a purchase (1,908). This imbalance suggests that predictive models may need adjustments to account for the disproportionate classes.

4.3. Correlation and Co-Variation Analysis

To assess relationships among predictors and with the outcome, we analyzed correlations and conducted statistical tests. Numeric predictors showed a strong positive correlation between BounceRates and ExitRates ($r = 0.91$), suggesting a potential overlap in what they measure, which could complicate modeling. ProductRelated_Duration had weak correlations with other variables, indicating it may offer unique predictive power. Tests of association revealed significant relationships between Revenue and VisitorType, Weekend, and Month (all p -values < 0.002), confirming their relevance as predictors. Additionally, ProductRelated_Duration varied significantly between buyers and non-buyers ($p < 0.001$), reinforcing its importance.

4.4. Data Pre-Processing

Based on our descriptive analytics, we examined several preprocessing steps and transformations to prepare the data for predictive modeling. To assess potential multicollinearity, we inspected the correlation matrix of numeric predictors and found a strong positive correlation between BounceRates and ExitRates ($r = 0.91$,). This high correlation suggests these variables may measure overlapping aspects of session abandonment, which could destabilize our regression models by inflating variance. To address this, we plan to either select one variable (e.g., ExitRates, due to its broader scope) or create a composite metric, ensuring cleaner model interpretation.

We also analyzed the distribution of the outcome variable, Revenue, to understand its balance and aid in designing our models. The bar plot revealed a significant imbalance, with 84.5% non-purchases (10,422 sessions) and 15.5% purchases (1,908 sessions). This skewness indicates that a standard model might overpredict non-purchases, so we are considering oversampling the minority class (purchases) or adjusting the classification threshold to improve sensitivity to the positive class, which is critical for our goal of identifying purchase drivers.

Finally, to evaluate predictor distributions for modeling assumptions, we inspected the summary statistics and found that ProductRelated_Duration exhibits high variability ($SD = 1,913.67$ seconds) and a skewed distribution (e.g., $\max = 63,973.52$ vs. $\text{median} = 598.94$). This analysis suggests that a log transformation could normalize this variable, reducing the influence of extreme values and improving the linearity assumption in our Logistic Regression model. These steps—addressing multicollinearity, class imbalance, and skewness—stem from our descriptive findings and are designed to enhance the robustness and interpretability of our predictive models.

5. Modeling Methods and Model Specifications

5.1. Initial Model Specification

Our initial model predicts Revenue (binary) using ProductRelated_Duration, BounceRates, ExitRates, VisitorType, Weekend, and Month. These predictors reflect key aspects of online shopping: time on product pages measures engagement, bounce and exit rates capture abandonment, visitor type indicates loyalty, weekend visits reflect browsing context, and month accounts for seasonal effects. Based on its skewness, we are considering a log transformation for ProductRelated_Duration.

5.2. Initial Logistic Regression Model Results

Given the binary nature of Revenue, we fitted a preliminary Logistic Regression model with the specified predictors. The results show that ProductRelated_Duration, ExitRates, and VisitorType significantly influence purchase likelihood ($p < 0.001$), aligning with our descriptive analysis. Specifically, longer time on product pages and lower exit rates increase the odds of a purchase, while returning visitors are less likely to buy compared to new visitors. Certain months also matter: November boosts purchases ($p < 0.001$), while February, March, and May reduce them ($p < 0.05$). BounceRates and Weekend showed no significant effects ($p > 0.4$), suggesting limited independent impact after accounting for other predictors.

5.3. Assumption Tests

To assess multicollinearity among predictors in our logistic regression model, we calculated Variance Inflation Factors (VIF). All GVIF-adjusted values were well below the common threshold of 5, indicating no significant multicollinearity. Specifically, BounceRates and ExitRates had VIF values of 1.75 and 1.79, respectively, suggesting moderate correlation but not at a level that compromises model stability. All other predictors, including ProductRelated_Duration, VisitorType, Weekend, and Month, had GVIF values near 1. As such, we conclude that multicollinearity is not a concern in our logistic regression model, and all predictors were retained for interpretability and completeness.

5.4. Model Candidates and Rationale

The first model is a standard logistic regression, which is appropriate for the binary outcome variable (Revenue). Outside of this control model, we evaluated two alternative modeling methods to address challenges that came up during our descriptive analysis.

For the second model, we selected LASSO Regression to maximize model simplicity and reduce potential redundancy among the predictors. Although our tests for multicollinearity (i.e., VIF) did not indicate severe multicollinearity, we did observe a high correlation ($r=0.91$) between BounceRates and ExitRates. LASSO applies a penalty that can shrink small coefficients to zero, effectively removing less informative predictors. This helps to identify which session activities and user attributes are most important for predicting purchase behavior.

The last model we chose is an Ensemble Tree-based approach that covers Bagging and Random Forest. These methods are useful for our binary outcome, handle complex interactions, and non-linearities without prior transformations.

5.5. Model Specification Candidates and Rationale

The first candidate model selected was a logistic regression using the full set of predictors related to the business case (ProductRelated_Duration, BounceRates, ExitRates, VisitorType, and Weekend). This model reflects our business understanding that the probability of a purchase being made is based on session behavior (e.g., time spent on product pages, bounce rates, exit rates) and visitor attributes (e.g., returning customers, time of visits). The second specification for the logistic regression uses the same set of predictors but fits the model in a balanced version of the dataset, where the number of purchase sessions and non-purchase sessions was equalized through random oversampling of the minority class (sessions which resulted in a purchase). Oversampling reduces the extreme skew of data (15.5% purchases) without eliminating any non-purchase data. This approach allows fairer evaluation of purchase drivers without the loss of useful behavior data.

The first LASSO model was fitted on the original, unbalanced dataset (Shopping), which had 15.5% purchase sessions. We used 10-fold cross-validation to identify the optimal regularization strength. The model selected a best lambda of 0.00041, which determines the amount of shrinkage applied to the coefficients. The coefficient path plot shows how coefficients shrink as the penalty increases, while the cross-validation plot illustrates how model deviance changes

with varying levels of regularization. To better evaluate purchase behavior without bias toward the majority class, we created a second LASSO model using the balanced version of the dataset (Shopping.bal), where the number of purchase and non-purchase sessions was equalized through oversampling. Cross-validation identified a best lambda of 0.00094 for this model. Coefficient path and CV error plots for this specification are shown in and, respectively.

For the first specification of the ensemble tree methods, we used Bagging with unbalanced data to obtain optimal prediction accuracy concerning the true class ratios. This process combines results from multiple trees that are built on bootstrap samples to limit variance and avoid overfitting. For the second specification, we used the balanced dataset to determine if equal class representation improves the prediction performance of the minority class. For the second specification, we used Random Forest with both datasets (unbalanced and balanced) with a low mtry value to introduce feature-level randomness to enable variance importance ranking.

5.6 Cross-Validation Testing and Final Model Selection

We evaluated all candidate models and specifications using 10-fold cross-validation (10FCV) to assess the model's predictive accuracy and select the most effective model. Each model was tested under two specifications: Spec 1 using the original, unbalanced dataset (15.5% purchase sessions), and Spec 2 using a balanced version created through oversampling. The results highlight Bagging with the balanced dataset as the best-performing model, achieving the lowest 10FCV error of 0.06592. This model outperformed all others, including Random Forest (Balanced) with a 10FCV error of 0.23570 and Logit Spec 1 (Unbalanced) at 0.11820. While logistic regression models offer interpretability, they were outperformed in predictive power by ensemble methods. LASSO, though useful for feature selection, produced higher CV errors in both specifications. Given our goal of maximizing predictive accuracy, we recommend the Bagging model trained on the balanced data as the final model. It delivers the strongest generalization performance and is best suited for deployment in a prediction-focused context.

The final two models were tree-based classifiers: Bagging and Random Forest. Both models were estimated using the same six predictors used in the regression-based models: ProductRelated_Duration, BounceRates, ExitRates, VisitorType, Weekend, and Month. For each approach, we fit a model on both the unbalanced dataset (Shopping) and the balanced version (Shopping.bal). The bagging model used all predictors at each split (mtry = 6), while the random forest model randomly selects a subset (mtry = 3), introducing additional variation between trees. These models were chosen because they handle nonlinear relationships and interactions between predictors without requiring explicit specification. We relied on the built-in classification error provided by the randomForest function, as it uses internal resampling to estimate performance. Comparing the balanced and unbalanced specifications allows us to assess the impact of class imbalance on prediction accuracy and the ability of tree-based methods to capture complex patterns in customer purchase behavior.

Final Method/Specification Selected:

While the bagging model with balanced data had the lowest 10-fold cross-validation error (0.0659), we prioritized interpretability to best serve our audience of e-commerce stakeholders.

Consequently, we selected the logistic regression model with the unbalanced specification, which produced a strong 10FCV error of 0.1182, making it the best-performing parametric model.

We fit this model on the entire unbalanced dataset using `glm()`, maintaining all core predictors from our initial specification. Although `BounceRates` was statistically insignificant and highly collinear with `ExitRates`, we retained it for interpretability and comparison across models. Its presence does not negatively impact the model's fit or generalization error. This final model provides a clear, interpretable view of the key behavioral factors driving online purchases and supports strategic business decision-making for conversion optimization.

6. Analysis of Results

The logistic regression model (unbalanced specification) was ultimately selected for our final analysis. Although the balanced bagging model had the lowest 10-FCV error overall (0.0659), we prioritized balancing interpretability and predictive accuracy for business stakeholders. This logistic model yielded a strong 10-fold cross-validation error of 0.1182, and the output displayed all core predictors as statistically significant, except `BounceRates`, which we retained for comparison. The model was fit using `glm()` on the full unbalanced dataset.

Quantitative predictors: All interpretations are “on average and holding everything else constant”:

A one-second increase in `ProductRelated_Duration` increases the odds of purchase by 0.0085% (coefficient = 0.000085, $p < 0.001$). This confirms that shoppers who spend more time engaging with product pages are more likely to convert.

A one-unit increase in `ExitRates` decreases the odds of purchase by nearly 100% (coefficient = -31.72, $p < 0.001$). This is the strongest negative predictor in the model and highlights that high exit behavior significantly reduces conversion probability.

Binary predictors: All interpretations are “on average and holding everything else constant”:

Returning Visitors have 30% lower odds of making a purchase compared to new visitors (coefficient = -0.353, $p < 0.001$). This suggests that returning users, despite their familiarity, are less likely to convert during a session.

`VisitorTypeOther` had a positive effect (coefficient = 0.295), but it was not statistically significant ($p = 0.342$).

Weekend visits slightly increase the odds of purchase by 4.7% (coefficient = 0.046), but this effect is also not statistically significant ($p = 0.443$).

Categorical predictors (Month): Relative to January (baseline month):

November significantly increases the odds of purchase by 60% (coefficient = 0.469, $p < 0.001$), aligning with expected seasonal shopping behavior.

February, March, and May all significantly decrease the odds of purchase:

February: -87% (coefficient = -2.034, $p = 0.0007$)

March: -42% (coefficient = -0.552, $p = 0.0003$)
May: -30% (coefficient = -0.359, $p = 0.0125$)

Other months, including July, June, October, and September, did not show significant differences in purchase likelihood relative to January.

7. Conclusions and Lessons Learned

7.1. *Conclusions from the Analysis*

The central business question in this project was: *Can we predict which online shopping sessions are likely to result in a purchase based on session behavior and visitor characteristics?* All models identified ProductRelated_Duration and ExitRates as the most consistent and influential predictors. Sessions with longer product interaction time and lower exit rates were more likely to lead to purchases. BounceRates also contributed to predictive performance, but to a lesser extent.

We developed and evaluated four classification models: logistic regression, LASSO regression, Bagging, and Random Forest. The initial logistic regression model (unbalanced) yielded a 10-fold cross-validation (10FCV) error rate of 11.8%, while the balanced version increased slightly to 20.2%. LASSO performed worse, with 10FCV error rates of 75.7% (unbalanced) and 116.9% (balanced), indicating it was not well-suited to this problem given the predictor structure and data distribution.

Tree-based models outperformed both regression approaches. Bagging achieved a 10FCV error of 17.6% on the unbalanced data and improved to 6.6% on the balanced data. Random Forest produced a 10FCV error of 15.5% unbalanced and 23.6% balanced. These models also ranked variable importance, reaffirming that product engagement metrics were far more predictive of purchases than categorical variables like VisitorType or Month. Unlike regression models, which assume linear relationships and require manual interaction terms, tree models handle nonlinearity and interaction effects automatically, resulting in stronger predictive performance.

In sum, if the goal is to understand general drivers of purchase behavior, logistic regression offers a solid balance of interpretability and performance. However, if the goal is to accurately classify purchase sessions for operational targeting, Bagging on the balanced dataset delivered the lowest overall error and the most reliable performance.

Based on our results, we recommend that online retailers target user engagement through optimization of high-exit pages and schedule targeted ad campaigns for repeat buyers. Seasonal campaigns, especially around November (the holiday season), need to be given priority to take advantage of higher conversion potential. Implementing immediate, automated actions during the duration of user sessions can also reduce exit rates and increase overall purchases made.

7.2. *Project Issues, Challenges and Lessons Learned*

A primary challenge in this project was the significant class imbalance, with only 15.5% of sessions resulting in a purchase. Models trained on the original dataset tended to favor the

majority class, reducing their ability to correctly identify purchasing sessions. This issue was addressed through oversampling, which allowed for more balanced learning and substantially improved classification performance across models.

In addition to outcome imbalance, we encountered limitations within the predictor set. BounceRates and ExitRates were found to be highly correlated, suggesting potential redundancy. While this had minimal impact on tree-based models, it may have introduced multicollinearity concerns in regression-based models, affecting coefficient interpretation and stability. Furthermore, ProductRelated_Duration exhibited a highly skewed distribution with several extreme values. These outliers may have influenced model training and contributed to differences in performance across methods.

The relative performance of the models varied by technique. LASSO facilitated variable selection but underperformed in terms of predictive accuracy. Logistic regression offered interpretability and performed competitively on the unbalanced data. Bagging and Random Forest achieved the highest predictive accuracy, particularly when trained on the balanced dataset. Additionally, error estimates obtained through 10-fold cross-validation were consistent with internal validation metrics for ensemble models, reinforcing the reliability of our evaluation approach.

This project emphasized the importance of addressing outcome imbalance, carefully assessing predictor redundancy and distribution, and selecting modeling approaches that align with the objectives of either explanation or prediction.

Appendix Contents

Table of contents

4.1 Descriptive statistics of key variables.....	1
4.2 Distribution of the outcome variable (e.g., normal, skewed, Poisson, binary, etc.- BINARY	6
4.3 Correlation and co-variation analysis.....	8
Chi-square Test (cat v. cat)	9
ANOVA Test (num vs. cat)-product duration differ by Revenue.....	10
5.2 Initial Logistic Regression Model.....	10
5.5 Model Candidates Specification Logit model with balanced data	12
LASSO Model Unbalanced Specification.....	13
LASSO Model Balanced Specification.....	15
Bagging Model	17
Random Forest	20
5.6 - 10FCV For Logit, LASSO ,Random Forest, and Bagging models.....	23
References.....	24

Load Dataset

```
Shopping <- read.table("online_shoppers_intention.csv", header = TRUE, sep =  
",")
```

4.1 Descriptive statistics of key variables

```
library(dplyr)
```

```
# Convert categorical variables to factors  
Shopping$Month <- as.factor(Shopping$Month)  
Shopping$VisitorType <- as.factor(Shopping$VisitorType)  
Shopping$OperatingSystems <- as.factor(Shopping$OperatingSystems)  
Shopping$Browser <- as.factor(Shopping$Browser)  
Shopping$Region <- as.factor(Shopping$Region)  
Shopping$TrafficType <- as.factor(Shopping$TrafficType)
```

```
head(Shopping)
```

	Administrative	Administrative_Duration	Informational	Informational_Duration
1	0	0	0	0
2	0	0	0	0

3	0	0	0	0
4	0	0	0	0
5	0	0	0	0
6	0	0	0	0

	ProductRelated	ProductRelated_Duration	BounceRates	ExitRates	PageValues
1	1	0.000000	0.20000000	0.2000000	0
2	2	64.000000	0.00000000	0.1000000	0
3	1	0.000000	0.20000000	0.2000000	0
4	2	2.666667	0.05000000	0.1400000	0
5	10	627.500000	0.02000000	0.0500000	0
6	19	154.216667	0.01578947	0.0245614	0

	SpecialDay	Month	OperatingSystems	Browser	Region	TrafficType
1	0	Feb	1	1	1	1
2	0	Feb	2	2	1	2
3	0	Feb	4	1	9	3
4	0	Feb	3	2	2	4
5	0	Feb	3	3	1	4
6	0	Feb	2	2	1	3

	VisitorType	Weekend	Revenue
1	Returning_Visitor	FALSE	FALSE
2	Returning_Visitor	FALSE	FALSE
3	Returning_Visitor	FALSE	FALSE
4	Returning_Visitor	FALSE	FALSE
5	Returning_Visitor	TRUE	FALSE
6	Returning_Visitor	FALSE	FALSE

summary(Shopping)

Administrative	Administrative_Duration	Informational
Min. : 0.000	Min. : 0.00	Min. : 0.0000
1st Qu.: 0.000	1st Qu.: 0.00	1st Qu.: 0.0000
Median : 1.000	Median : 7.50	Median : 0.0000
Mean : 2.315	Mean : 80.82	Mean : 0.5036
3rd Qu.: 4.000	3rd Qu.: 93.26	3rd Qu.: 0.0000
Max. : 27.000	Max. : 3398.75	Max. : 24.0000

Informational_Duration	ProductRelated	ProductRelated_Duration
Min. : 0.00	Min. : 0.00	Min. : 0.0
1st Qu.: 0.00	1st Qu.: 7.00	1st Qu.: 184.1
Median : 0.00	Median : 18.00	Median : 598.9
Mean : 34.47	Mean : 31.73	Mean : 1194.8
3rd Qu.: 0.00	3rd Qu.: 38.00	3rd Qu.: 1464.2
Max. : 2549.38	Max. : 705.00	Max. : 63973.5

BounceRates	ExitRates	PageValues	SpecialDay
Min. : 0.000000	Min. : 0.00000	Min. : 0.000	Min. : 0.00000
1st Qu.: 0.000000	1st Qu.: 0.01429	1st Qu.: 0.000	1st Qu.: 0.00000
Median : 0.003112	Median : 0.02516	Median : 0.000	Median : 0.00000
Mean : 0.022191	Mean : 0.04307	Mean : 5.889	Mean : 0.06143
3rd Qu.: 0.016813	3rd Qu.: 0.05000	3rd Qu.: 0.000	3rd Qu.: 0.00000

Max. :0.200000 Max. :0.200000 Max. :361.764 Max. :1.000000

Month	OperatingSystems	Browser	Region	TrafficType
May :3364	2 :6601	2 :7961	1 :4780	2 :3913
Nov :2998	1 :2585	1 :2462	3 :2403	1 :2451
Mar :1907	3 :2555	4 : 736	4 :1182	3 :2052
Dec :1727	4 : 478	5 : 467	2 :1136	4 :1069
Oct : 549	8 : 79	6 : 174	6 : 805	13 : 738
Sep : 448	6 : 19	10 : 163	7 : 761	10 : 450
(Other):1337	(Other): 13	(Other): 367	(Other):1263	(Other):1657

VisitorType	Weekend	Revenue
New_Visitor : 1694	Mode :logical	Mode :logical
Other : 85	FALSE:9462	FALSE:10422
Returning_Visitor:10551	TRUE :2868	TRUE :1908

Mean and Standard Dev. for Numeric Variables

```
library(psych)
describe(Shopping)
```

	vars	n	mean	sd	median	trimmed	mad	min
Administrative	1	12330	2.32	3.32	1.00	1.63	1.48	0
Administrative_Duration	2	12330	80.82	176.78	7.50	42.10	11.12	0
Informational	3	12330	0.50	1.27	0.00	0.18	0.00	0
Informational_Duration	4	12330	34.47	140.75	0.00	3.59	0.00	0
ProductRelated	5	12330	31.73	44.48	18.00	22.75	19.27	0
ProductRelated_Duration	6	12330	1194.75	1913.67	598.94	820.08	742.69	0
BounceRates	7	12330	0.02	0.05	0.00	0.01	0.00	0
ExitRates	8	12330	0.04	0.05	0.03	0.03	0.02	0
PageValues	9	12330	5.89	18.57	0.00	1.29	0.00	0
SpecialDay	10	12330	0.06	0.20	0.00	0.00	0.00	0
Month*	11	12330	6.16	2.37	7.00	6.35	1.48	1
OperatingSystems*	12	12330	2.12	0.91	2.00	2.06	0.00	1
Browser*	13	12330	2.36	1.72	2.00	2.00	0.00	1
Region*	14	12330	3.15	2.40	3.00	2.79	2.97	1
TrafficType*	15	12330	4.07	4.03	2.00	3.21	1.48	1
VisitorType*	16	12330	2.72	0.69	3.00	2.90	0.00	1
Weekend	17	12330	NaN	NA	NA	NaN	NA	Inf
Revenue	18	12330	NaN	NA	NA	NaN	NA	Inf

	max	range	skew	kurtosis	se
Administrative	27.00	27.00	1.96	4.70	0.03
Administrative_Duration	3398.75	3398.75	5.61	50.53	1.59
Informational	24.00	24.00	4.04	26.92	0.01
Informational_Duration	2549.38	2549.38	7.58	76.27	1.27
ProductRelated	705.00	705.00	4.34	31.19	0.40
ProductRelated_Duration	63973.52	63973.52	7.26	137.10	17.23

BounceRates	0.20	0.20	2.95	7.72	0.00
ExitRates	0.20	0.20	2.15	4.01	0.00
PageValues	361.76	361.76	6.38	65.60	0.17
SpecialDay	1.00	1.00	3.30	9.91	0.00
Month*	10.00	9.00	-0.83	-0.37	0.02
OperatingSystems*	8.00	7.00	2.07	10.45	0.01
Browser*	13.00	12.00	3.24	12.74	0.02
Region*	9.00	8.00	0.98	-0.15	0.02
TrafficType*	20.00	19.00	1.96	3.48	0.04
VisitorType*	3.00	2.00	-2.06	2.29	0.01
Weekend	-Inf	-Inf	NA	NA	NA
Revenue	-Inf	-Inf	NA	NA	NA

Mean Product Duration by Revenue

```
tapply(Shopping$ProductRelated_Duration, Shopping$Revenue, mean, na.rm = TRUE)
```

FALSE	TRUE
1069.988	1876.210

```
tapply(Shopping$ProductRelated_Duration, Shopping$Revenue, sd, na.rm = TRUE)
```

FALSE	TRUE
1803.798	2312.214

Bounce Rates by Revenue

```
tapply(Shopping$BounceRates, Shopping$Revenue, mean, na.rm = TRUE)
```

FALSE	TRUE
0.025317232	0.005117153

```
tapply(Shopping$BounceRates, Shopping$Revenue, sd, na.rm = TRUE)
```

FALSE	TRUE
0.05187746	0.01218451

Exit Rates by Revenue

```
tapply(Shopping$ExitRates, Shopping$Revenue, mean, na.rm = TRUE)
```

FALSE	TRUE
0.04737827	0.01955517

Page Values by Revenue

```
tapply(Shopping$PageValues, Shopping$Revenue, mean, na.rm = TRUE)
```

FALSE	TRUE
1.975998	27.264518

Special Day Score by Revenue

```
tapply(Shopping$SpecialDay, Shopping$Revenue, mean, na.rm = TRUE)
```

FALSE	TRUE
0.06843216	0.02316562

Categorical Variables Frequency

```
#Visitor Type
```

```
table(Shopping$VisitorType)
```

New_Visitor	Other	Returning_Visitor
1694	85	10551

```
table(Shopping$VisitorType, Shopping$Revenue)
```

	FALSE	TRUE
New_Visitor	1272	422
Other	69	16
Returning_Visitor	9081	1470

```
# Weekend visits
```

```
table(Shopping$Weekend)
```

FALSE	TRUE
9462	2868

```
table(Shopping$Weekend, Shopping$Revenue)
```

	FALSE	TRUE
FALSE	8053	1409
TRUE	2369	499

```
# Month of visit
```

```
table(Shopping$Month)
```

Aug	Dec	Feb	Jul	June	Mar	May	Nov	Oct	Sep
433	1727	184	432	288	1907	3364	2998	549	448

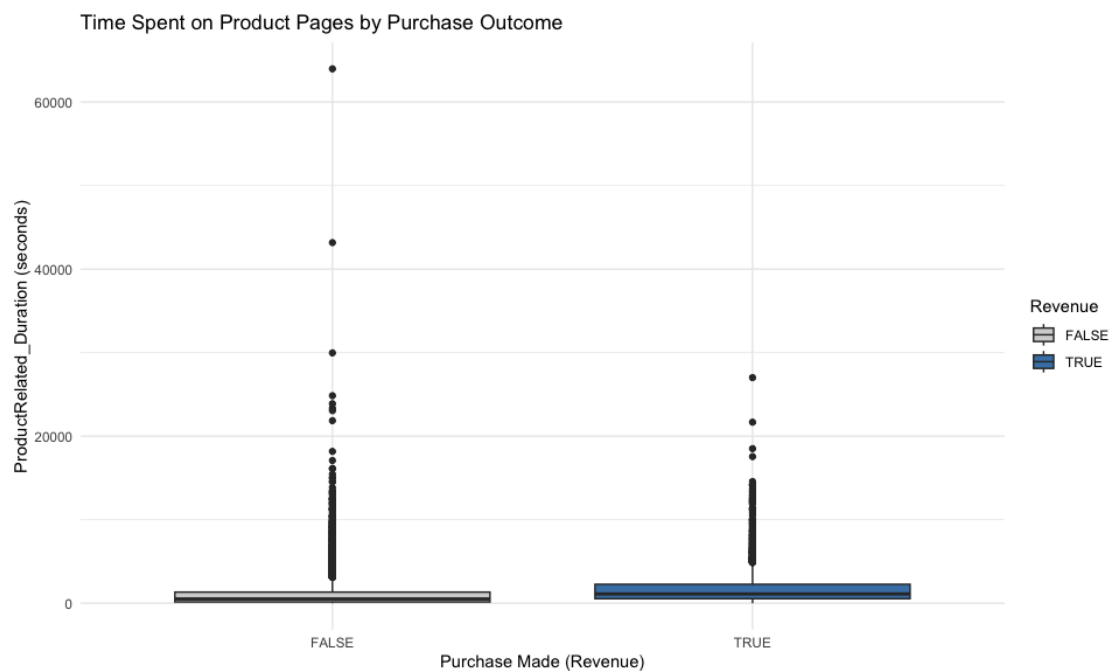
```
table(Shopping$Month, Shopping$Revenue)
```

	FALSE	TRUE
Aug	357	76
Dec	1511	216
Feb	181	3
Jul	366	66
June	259	29
Mar	1715	192
May	2999	365
Nov	2238	760
Oct	434	115
Sep	362	86

Boxplot for Time Spent on Product Pages by Purchase Outcome

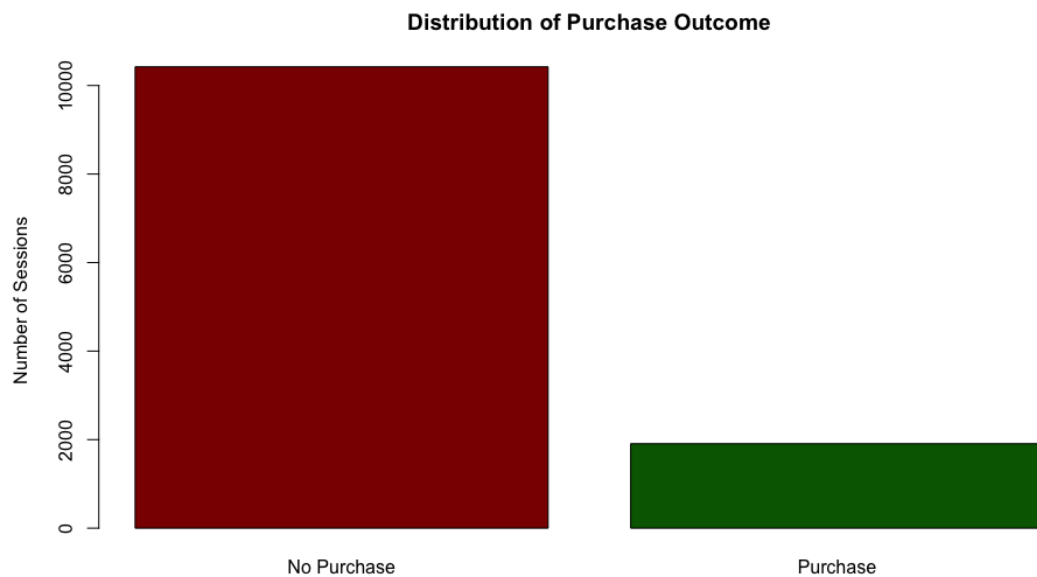
```
library(ggplot2)

# Boxplot of time spent on product pages by purchase outcome
ggplot(Shopping, aes(x = Revenue, y = ProductRelated_Duration, fill =
Revenue)) +
  geom_boxplot() +
  labs(title = "Time Spent on Product Pages by Purchase Outcome",
       x = "Purchase Made (Revenue)",
       y = "ProductRelated_Duration (seconds)") +
  scale_fill_manual(values = c("FALSE" = "lightgray", "TRUE" = "steelblue"))
+
  theme_minimal()
```



4.2 Distribution of the outcome variable (e.g., normal, skewed, Poisson, binary, etc.)-BINARY

```
barplot(table(Shopping$Revenue),
       main = "Distribution of Purchase Outcome",
       ylab = "Number of Sessions",
       names.arg = c("No Purchase", "Purchase"),
       col = c("darkred", "darkgreen"))
```



Correlation Heatmap of Numeric Variables

#Previous Code for simple heatmap

```
numeric_vars <- Shopping[, c("ProductRelated_Duration", "BounceRates",
                              "ExitRates", "PageValues", "SpecialDay")]
```

Correlation matrix

```
cor_matrix <- cor(numeric_vars, use = "complete.obs")
```

```
print(cor_matrix)
```

	ProductRelated_Duration	BounceRates	ExitRates
ProductRelated_Duration	1.00000000	-0.18454112	-0.2519841
BounceRates	-0.18454112	1.00000000	0.9130044
ExitRates	-0.25198410	0.91300440	1.00000000
PageValues	0.05282306	-0.11938603	-0.1744983
SpecialDay	-0.03637985	0.07270225	0.1022418

	PageValues	SpecialDay
ProductRelated_Duration	0.05282306	-0.03637985
BounceRates	-0.11938603	0.07270225
ExitRates	-0.17449831	0.10224180
PageValues	1.00000000	-0.06354127
SpecialDay	-0.06354127	1.00000000

```
library(pheatmap)
```

Simple heatmap

```
pheatmap(cor_matrix,
```

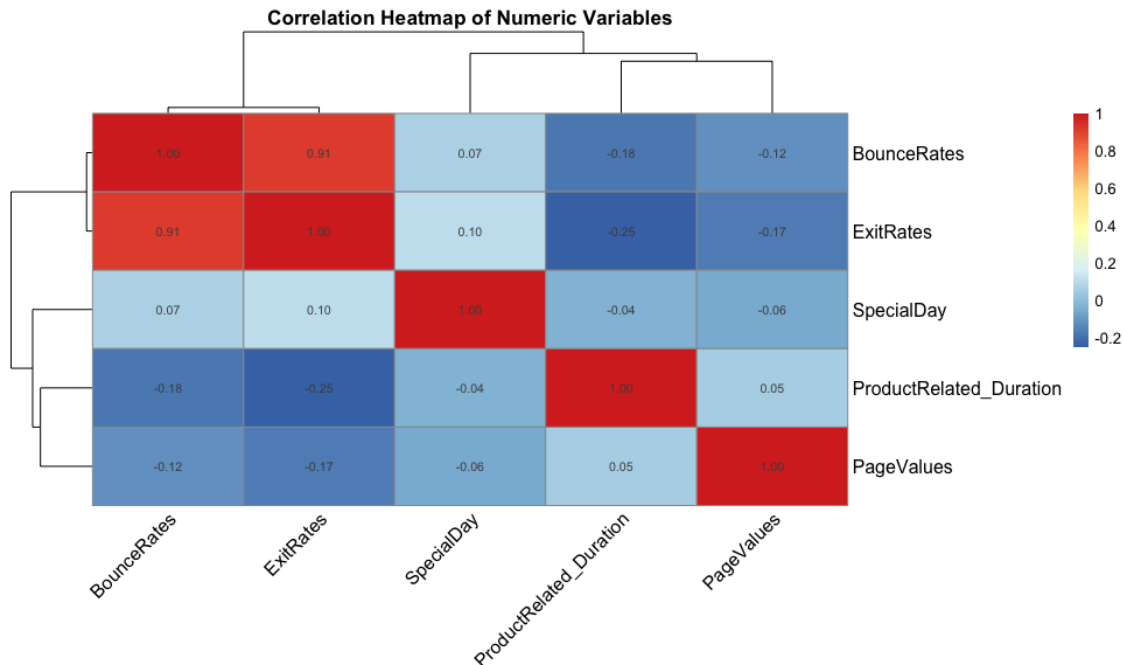
```
        main = "Correlation Heatmap of Numeric Variables",
```



```

fontsize_row = 12,      # Increase row label size
fontsize_col = 12,      # Increase column label size
display_numbers = TRUE, # Optional: show correlation values
clustering_method = "complete",
angle_col = 45)

```




4.3 Correlation and co-variation analysis

```

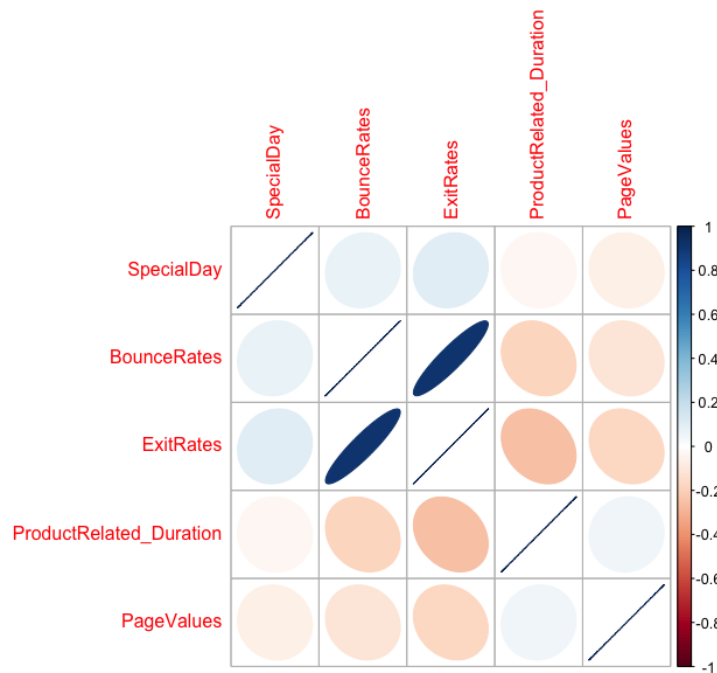
# numeric variables
Shopping.num <- Shopping[, c("ProductRelated_Duration", "BounceRates",
"ExitRates", "PageValues", "SpecialDay")]

# Correlation plots
Shopping.cor <- cor(Shopping.num)
library(corrplot)
corrplot(Shopping.cor,
order= "hclust",
method= "number")

```



```
corrplot(Shopping.cor,
order= "hclust",
method= "ellipse")
```



Chi-square Test (cat v. cat)

```
# Visitor Type vs Revenue
table_vt <- table(Shopping$VisitorType, Shopping$Revenue)
chisq.test(table_vt)
```

Pearson's Chi-squared test

```
data: table_vt
X-squared = 135.25, df = 2, p-value < 2.2e-16

# Weekend vs Revenue
table_weekend <- table(Shopping$Weekend, Shopping$Revenue)
chisq.test(table_weekend)
```

Pearson's Chi-squared test with Yates' continuity correction

```
data: table_weekend
X-squared = 10.391, df = 1, p-value = 0.001266

# Month vs Revenue
table_month <- table(Shopping$Month, Shopping$Revenue)
chisq.test(table_month)
```

Pearson's Chi-squared test

```
data: table_month
X-squared = 384.93, df = 9, p-value < 2.2e-16
```

ANOVA Test (num vs. cat)-product duration differ by Revenue

```
anova_result <- aov(ProductRelated_Duration ~ Revenue, data = Shopping)
summary(anova_result)
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
Revenue	1	1.048e+09	1.048e+09	293	<2e-16 ***
Residuals	12328	4.410e+10	3.577e+06		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

5.2 Initial Logistic Regression Model

```
# Fit the model
```

```
logit_model <- glm(Revenue ~ ProductRelated_Duration + BounceRates +  
ExitRates + VisitorType + Weekend + Month,  
family = binomial, data = Shopping)
```

```
# Summary of results  
summary(logit_model)
```

```
Call:  
glm(formula = Revenue ~ ProductRelated_Duration + BounceRates +
```

```
ExitRates + VisitorType + Weekend + Month, family = binomial,
data = Shopping)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	-5.906e-01	1.442e-01	-4.095	4.23e-05	***
ProductRelated_Duration	8.474e-05	1.227e-05	6.906	4.98e-12	***
BounceRates	8.528e-01	3.281e+00	0.260	0.794943	
ExitRates	-3.172e+01	2.173e+00	-14.597	< 2e-16	***
VisitorTypeOther	2.953e-01	3.111e-01	0.949	0.342439	
VisitorTypeReturning_Visitor	-3.525e-01	6.929e-02	-5.088	3.62e-07	***
WeekendTRUE	4.624e-02	6.033e-02	0.766	0.443434	
MonthDec	-3.746e-01	1.520e-01	-2.464	0.013729	*
MonthFeb	-2.034e+00	6.012e-01	-3.383	0.000718	***
MonthJul	6.023e-03	1.915e-01	0.031	0.974911	
MonthJune	-3.482e-01	2.427e-01	-1.435	0.151416	
MonthMar	-5.523e-01	1.532e-01	-3.606	0.000310	***
MonthMay	-3.585e-01	1.435e-01	-2.499	0.012459	*
MonthNov	4.693e-01	1.390e-01	3.377	0.000732	***
MonthOct	1.170e-01	1.704e-01	0.686	0.492472	
MonthSep	1.667e-02	1.806e-01	0.092	0.926457	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 10624.8 on 12329 degrees of freedom
Residual deviance: 9302.4 on 12314 degrees of freedom
AIC: 9334.4

Number of Fisher Scoring iterations: 7

Multicollinearity Test (VIF)

```
library(car)
vif(logit_model)
```

	GVIF	Df	GVIF^(1/(2*Df))
ProductRelated_Duration	1.111288	1	1.054176
BounceRates	1.753683	1	1.324267
ExitRates	1.798105	1	1.340934
VisitorType	1.166589	2	1.039273
Weekend	1.009531	1	1.004754
Month	1.096168	9	1.005114

Balancing the Data

```
set.seed(1) # Random seed
```

Separate the positive and negative classes

```
Shopping.1 <- subset(Shopping, Revenue == TRUE) # Purchases
```

```

Shopping.0 <- subset(Shopping, Revenue == FALSE) # Non-purchases

cat("Positive observations = ", nrow(Shopping.1), "\n")
Positive observations = 1908

cat("Negative observations = ", nrow(Shopping.0), "\n")
Negative observations = 10422

# Oversample positives to match the number of negatives
Shopping.1.oversampled <- Shopping.1[sample(nrow(Shopping.1), size =
nrow(Shopping.0), replace = TRUE), ]
cat("Oversampled positive observations = ", nrow(Shopping.1.oversampled),
"\n")

Oversampled positive observations = 10422

# Combine oversampled positives with original negatives
Shopping.bal <- rbind(Shopping.0, Shopping.1.oversampled)
cat("Total balanced observations = ", nrow(Shopping.bal), "\n")

Total balanced observations = 20844

# Check new class balance
table(Shopping.bal$Revenue)

FALSE  TRUE
10422 10422

```

5.5 Model Candidates Specification Logit model with balanced data

```

logit_model_bal <- glm(Revenue ~ ProductRelated_Duration + BounceRates +
ExitRates + VisitorType + Weekend + Month,
                      family = binomial, data = Shopping.bal)

```

```

# Summary of results
summary(logit_model_bal)

```

Call:

```

glm(formula = Revenue ~ ProductRelated_Duration + BounceRates +
    ExitRates + VisitorType + Weekend + Month, family = binomial,
    data = Shopping.bal)

```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	1.088e+00	8.773e-02	12.401	< 2e-16	***
ProductRelated_Duration	1.025e-04	9.102e-06	11.258	< 2e-16	***
BounceRates	2.355e+00	1.663e+00	1.416	0.15677	
ExitRates	-3.214e+01	1.176e+00	-27.329	< 2e-16	***

```

VisitorTypeOther      1.158e-02  2.030e-01  0.057  0.95451
VisitorTypeReturning_Visitor -4.259e-01  4.246e-02 -10.032 < 2e-16 ***
WeekendTRUE           2.583e-02  3.587e-02  0.720  0.47139
MonthDec              -3.263e-01  9.003e-02 -3.624  0.00029 ***
MonthFeb             -2.340e+00  3.280e-01 -7.134  9.73e-13 ***
MonthJul              9.160e-02  1.136e-01  0.807  0.41993
MonthJune            -3.279e-01  1.395e-01 -2.351  0.01872 *
MonthMar             -5.052e-01  9.004e-02 -5.611  2.02e-08 ***
MonthMay             -2.788e-01  8.529e-02 -3.268  0.00108 **
MonthNov              5.109e-01  8.485e-02  6.021  1.73e-09 ***
MonthOct              2.495e-01  1.038e-01  2.403  0.01626 *
MonthSep              1.471e-01  1.090e-01  1.350  0.17692

```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 28896 on 20843 degrees of freedom
Residual deviance: 24338 on 20828 degrees of freedom
AIC: 24370

Number of Fisher Scoring iterations: 6

Multicollinearity Check (again) for balanced data
vif(logit_model_bal)

	GVIF	Df	GVIF^(1/(2*Df))
ProductRelated_Duration	1.147739	1	1.071326
BounceRates	1.936491	1	1.391579
ExitRates	2.018948	1	1.420897
VisitorType	1.145355	2	1.034511
Weekend	1.009520	1	1.004749
Month	1.117384	9	1.006185

LASSO Model Unbalanced Specification

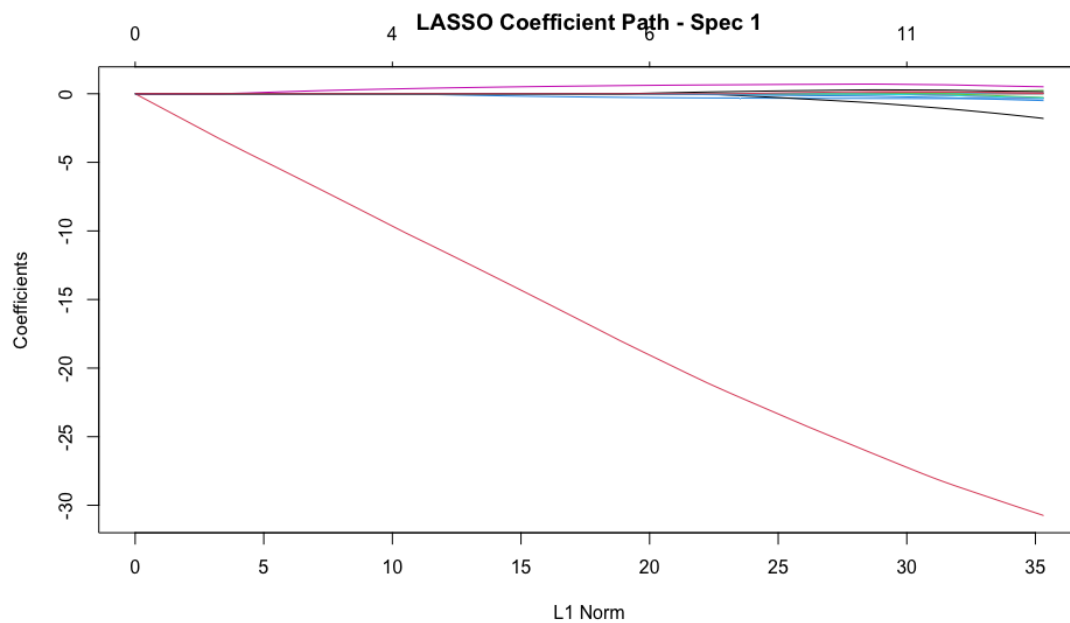
```

library(glmnet)


x1 <- model.matrix(Revenue ~ ProductRelated_Duration + BounceRates +
ExitRates +
                    VisitorType + Weekend + Month, data = Shopping)[, -1]
y1 <- Shopping$Revenue

# Fit LASSO model
lasso_model1 <- glmnet(x1, y1, alpha = 1, family = "binomial")
plot(lasso_model1, main = "LASSO Coefficient Path - Spec 1")

```



```
# Cross-validation
set.seed(1)
lasso_cv1 <- cv.glmnet(x1, y1, alpha = 1, family = "binomial")
plot(lasso_cv1)
```



```
# Extract optimal lambda and deviance
lasso_lambda1 <- lasso_cv1$lambda.min
lasso_coef1 <- coef(lasso_cv1, s = lasso_lambda1)
```

```
round(cbind("Best Lambda" = lasso_lambda1,
           "Best Log Lambda" = log(lasso_lambda1),
           "Best 10FCV" = min(lasso_cv1$cvm)), 5)
```

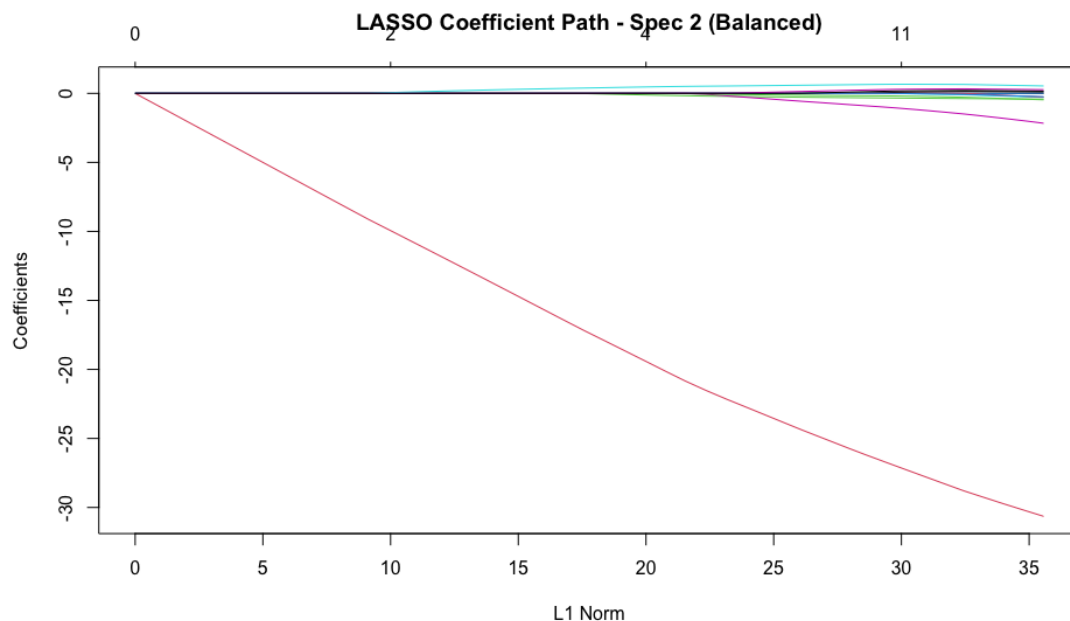
```
Best Lambda Best Log Lambda Best 10FCV
[1,]      0.00041      -7.80163      0.75661
```

LASSO Model Balanced Specification

```
library(glmnet)
```

```
# Create predictor matrix (x2) and outcome vector (y2) from balanced dataset
x2 <- model.matrix(Revenue ~ ProductRelated_Duration + BounceRates +
ExitRates +
                  VisitorType + Weekend + Month, data = Shopping.bal)[, -1]
y2 <- Shopping.bal$Revenue
```

```
# Fit LASSO model
lasso_model2 <- glmnet(x2, y2, alpha = 1, family = "binomial")
plot(lasso_model2, main = "LASSO Coefficient Path - Spec 2 (Balanced)")
```



```
# Cross-validation
set.seed(1)
lasso_cv2 <- cv.glmnet(x2, y2, alpha = 1, family = "binomial")
plot(lasso_cv2, main = "LASSO CV Error Plot - Spec 2 (Balanced)")
```



```
)
```

```
lasso.results
```

	Best Lambda	Best Log Lambda	Best 10FCV
Spec 1: Unbalanced	0.00041	-7.80163	0.75661
Spec 2: Balanced	0.00094	-6.96873	1.16945

10FCV for Both Logit Models

```
library(boot)
```

```
# Logistic Model - Spec 1: Unbalanced
```

```
set.seed(1)
```

```
cv_logit1 <- cv.glm(data = Shopping, glmfit = logit_model, K = 10)
```

```
# Logistic Model - Spec 2: Balanced
```

```
set.seed(1)
```

```
cv_logit2 <- cv.glm(data = Shopping.bal, glmfit = logit_model_bal, K = 10)
```

Bagging Model

```
# Set seed and RNG setting
```

```
RNGkind(sample.kind = "default")
```

```
set.seed(1)
```

```
# Load library
```

```
library(randomForest)
```

```
# Ensure binary outcome is a factor with labels
```

```
Shopping$Revenue <- factor(Shopping$Revenue, levels = c(FALSE, TRUE), labels  
= c("No", "Yes"))
```

```
Shopping.bal$Revenue <- factor(Shopping.bal$Revenue, levels = c(FALSE, TRUE),  
labels = c("No", "Yes"))
```

```
# Define modeling formula
```

```
model_formula <- Revenue ~ ProductRelated_Duration + BounceRates + ExitRates  
+
```

```
VisitorType + Weekend + Month
```

```
# Create NA-safe subsets for modeling
```

```
Shopping.clean <- na.omit(Shopping[, all.vars(model_formula)])
```

```
Shopping.bal.clean <- na.omit(Shopping.bal[, all.vars(model_formula)])
```

```
# -----
```

```
# 1. Bagging (Unbalanced, mtry = 6)
```

```
# -----
```

```
bag.unbal <- randomForest(model_formula,  
                           data = Shopping.clean,  
                           mtry = 6,
```

```
importance = TRUE)

print(bag.unbal)

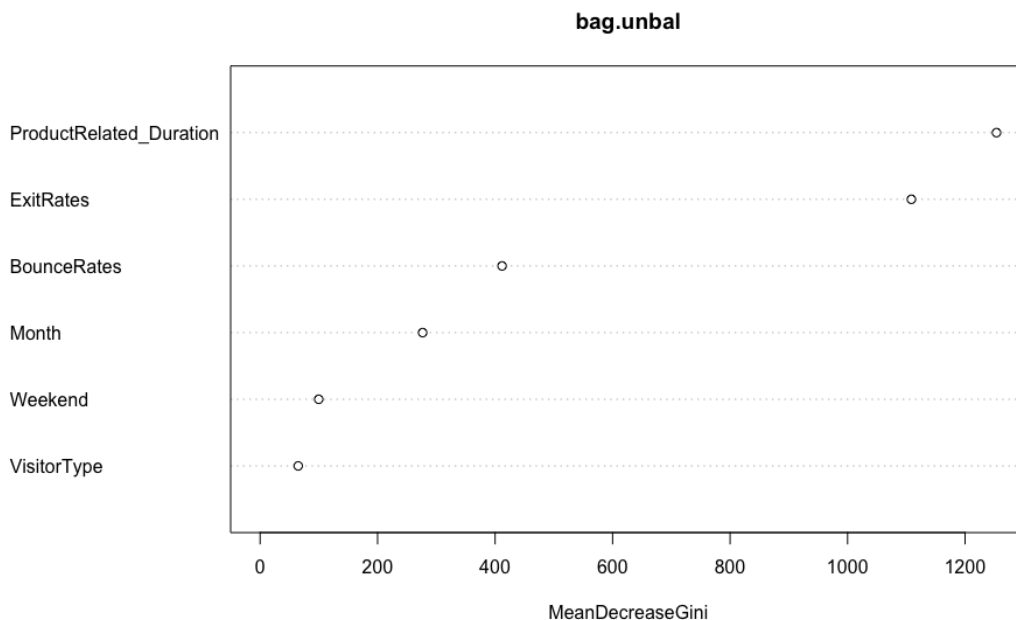
Call:
  randomForest(formula = model_formula, data = Shopping.clean,      mtry = 6,
importance = TRUE)
      Type of random forest: classification
      Number of trees: 500
No. of variables tried at each split: 6

      OOB estimate of  error rate: 16.94%
Confusion matrix:
      No Yes class.error
No  9951 471  0.04519286
Yes 1618 290  0.84800839

cat("\nClassification Error Rate (Bagging - Unbalanced) =",
    bag.unbal$err.rate[nrow(bag.unbal$err.rate), "OOB"])

Classification Error Rate (Bagging - Unbalanced) = 0.1694242

varImpPlot(bag.unbal, type = 2)
```



```
importance(bag.unbal, type = 2)

      MeanDecreaseGini
ProductRelated_Duration 1253.47359
```

```
BounceRates          411.85033
ExitRates             1108.58024
VisitorType           64.89726
Weekend               99.81519
Month                 276.65776
```

```
# -----
# 2. Bagging (Balanced, mtry = 6)
# -----
bag.bal <- randomForest(model_formula,
                        data = Shopping.bal.clean,
                        mtry = 6,
                        importance = TRUE)

print(bag.bal)
```


```
Call:
 randomForest(formula = model_formula, data = Shopping.bal.clean,      mtry =
6, importance = TRUE)
      Type of random forest: classification
      Number of trees: 500
No. of variables tried at each split: 6
```

```
      OOB estimate of  error rate: 5.16%
Confusion matrix:
      No  Yes class.error
No  9390  1032 0.099021301
Yes   43 10379 0.004125888
```

```
cat("\nClassification Error Rate (Bagging - Balanced) =",
    bag.bal$err.rate[nrow(bag.bal$err.rate), "OOB"])
```

```
Classification Error Rate (Bagging - Balanced) = 0.05157359
```

```
varImpPlot(bag.bal, type = 2)
```



```
importance(bag.bal, type = 2)
```

	MeanDecreaseGini
ProductRelated_Duration	3625.9984
BounceRates	1327.4517
ExitRates	3928.4428
VisitorType	188.3419
Weekend	260.4771
Month	1073.6137

Random Forest

```
# 3. Random Forest (Unbalanced, mtry = 3)
# -----
rf.unbal <- randomForest(model_formula,
                          data = Shopping.clean,
                          mtry = 3,
                          importance = TRUE)

print(rf.unbal)
```

Call:

```
randomForest(formula = model_formula, data = Shopping.clean, mtry = 3,
importance = TRUE)
```

Type of random forest: classification

Number of trees: 500

No. of variables tried at each split: 3

OOB estimate of error rate: 16.26%

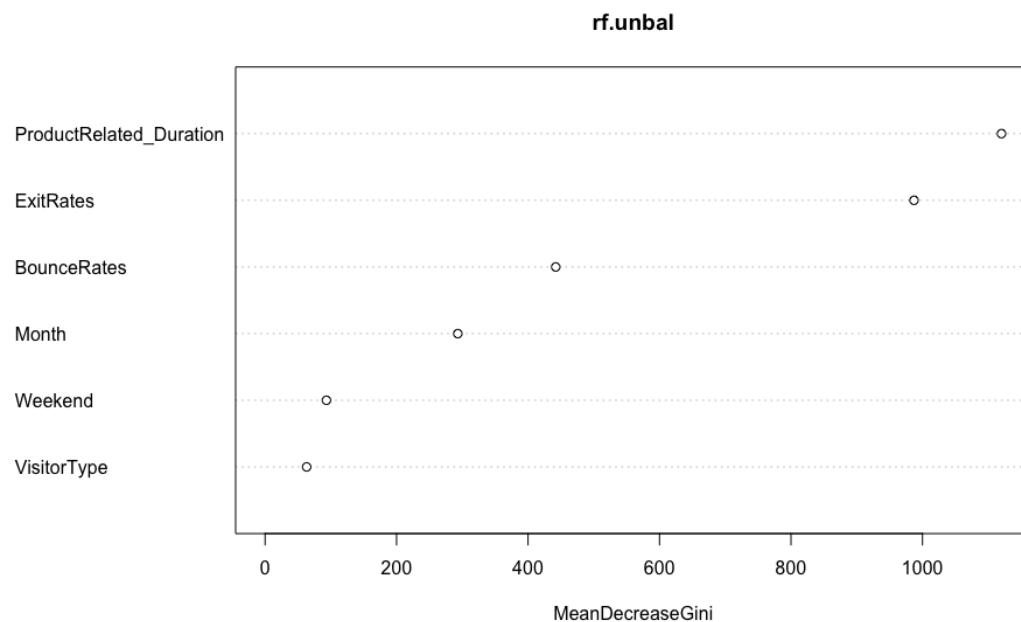
Confusion matrix:

	No	Yes	class.error
No	10056	366	0.03511802
Yes	1639	269	0.85901468

```
cat("\nClassification Error Rate (Random Forest - Unbalanced) =",  
    rf.unbal$err.rate[nrow(rf.unbal$err.rate), "OOB"])
```

Classification Error Rate (Random Forest - Unbalanced) = 0.1626115

```
varImpPlot(rf.unbal, type = 2)
```



```
importance(rf.unbal, type = 2)
```

	MeanDecreaseGini
ProductRelated_Duration	1120.22637
BounceRates	442.13433
ExitRates	987.14257
VisitorType	63.23252
Weekend	93.32027
Month	293.18836

```
# -----
```

```
# 4. Random Forest (Balanced, mtry = 3)
```

```
# -----
```

```
rf.bal <- randomForest(model_formula,  
                        data = Shopping.bal.clean,  
                        mtry = 3,  
                        importance = TRUE)
```

```
print(rf.bal)
```

Call:

```
randomForest(formula = model_formula, data = Shopping.bal.clean, mtry = 3, importance = TRUE)
```

Type of random forest: classification

Number of trees: 500

No. of variables tried at each split: 3

OOB estimate of error rate: 5.72%

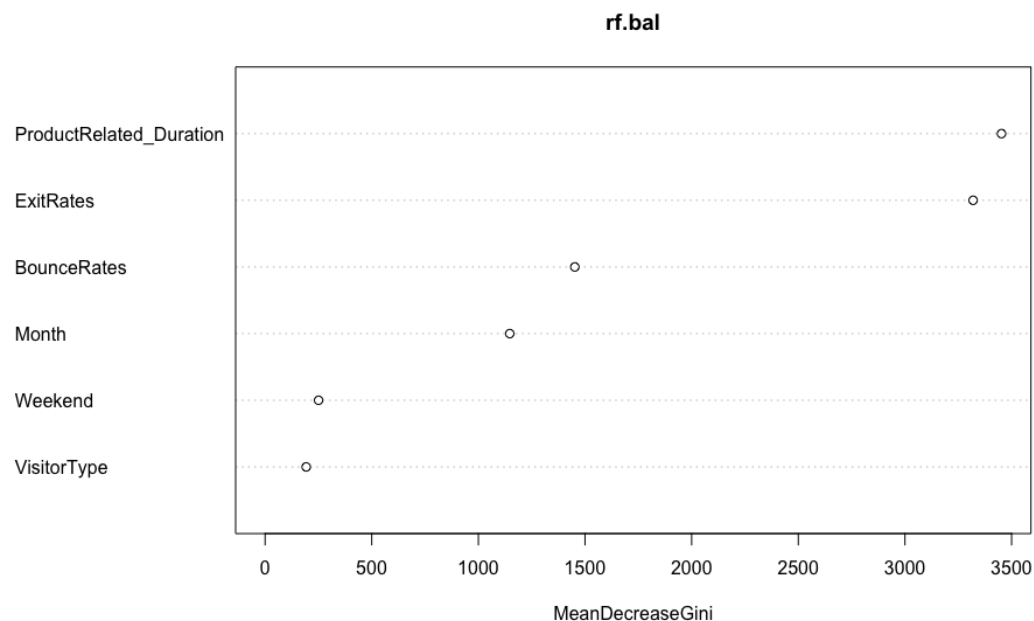
Confusion matrix:

	No	Yes	class.error
No	9274	1148	0.110151602
Yes	44	10378	0.004221838

```
cat("\nClassification Error Rate (Random Forest - Balanced) =",  
    rf.bal$err.rate[nrow(rf.bal$err.rate), "OOB"])
```

Classification Error Rate (Random Forest - Balanced) = 0.05718672

```
varImpPlot(rf.bal, type = 2)
```



```
importance(rf.bal, type = 2)
```

	MeanDecreaseGini
ProductRelated_Duration	3452.6982
BounceRates	1452.1695

ExitRates	3319.5724
VisitorType	192.7371
Weekend	250.7165
Month	1147.1118

5.6 - 10FCV For Logit, LASSO ,Random Forest, and Bagging models

10FCV for Bagging and Random Forest

```
# Load caret and randomForest
library(caret)
library(randomForest)

# Define formula (already established)
model_formula <- Revenue ~ ProductRelated_Duration + BounceRates + ExitRates
+ VisitorType + Weekend + Month

# Set control for 10-fold cross-validation
ctrl <- trainControl(method = "cv", number = 10)

# 1. Bagging - Unbalanced
set.seed(1)
bag_cv_unbal <- train(model_formula, data = Shopping.clean,
                      method = "treebag", trControl = ctrl)

# 2. Bagging - Balanced
set.seed(1)
bag_cv_bal <- train(model_formula, data = Shopping.bal.clean,
                    method = "treebag", trControl = ctrl)

# 3. Random Forest - Unbalanced
set.seed(1)
rf_cv_unbal <- train(model_formula, data = Shopping.clean,
                    method = "rf", trControl = ctrl, tuneGrid =
data.frame(mtry = 3))

# 4. Random Forest - Balanced
set.seed(1)
rf_cv_bal <- train(model_formula, data = Shopping.bal.clean,
                  method = "rf", trControl = ctrl, tuneGrid =
data.frame(mtry = 3))
```

Conversion for Bagging and Random Forest

```
# Extract and convert accuracy to error rate
bag_cv_unbal_error <- 1 - max(bag_cv_unbal$results$Accuracy)
bag_cv_bal_error <- 1 - max(bag_cv_bal$results$Accuracy)
rf_cv_unbal_error <- 1 - max(rf_cv_unbal$results$Accuracy)
rf_cv_bal_error <- 1 - max(rf_cv_bal$results$Accuracy)
```


Lasso and Logit Assignment for 10FCV

```
lasso_cv_error1 <- min(lasso_cv1$cvm)
lasso_cv_error2 <- min(lasso_cv2$cvm)
logit_cv_error1 <- cv_logit1$delta[2]
logit_cv_error2 <- cv_logit2$delta[2]
```

Complete Model Comparison of 10FCV Errors

```
model_comparison <- data.frame(
  Model = c("Logit Spec 1 (Unbalanced)", "Logit Spec 2 (Balanced)",
            "LASSO Spec 1 (Unbalanced)", "LASSO Spec 2 (Balanced)",
            "Bagging - Unbalanced", "Bagging - Balanced",
            "Random Forest - Unbalanced", "Random Forest - Balanced"),
  `10FCV Error` = round(c(logit_cv_error1, logit_cv_error2,
                          lasso_cv_error1, lasso_cv_error2,
                          bag_cv_unbal_error, bag_cv_bal_error,
                          rf_cv_unbal_error, rf_cv_bal_error), 5)
)
```

View

model_comparison

	Model	X10FCV.Error
1	Logit Spec 1 (Unbalanced)	0.11820
2	Logit Spec 2 (Balanced)	0.20225
3	LASSO Spec 1 (Unbalanced)	0.75661
4	LASSO Spec 2 (Balanced)	1.16945
5	Bagging - Unbalanced	0.17583
6	Bagging - Balanced	0.06592
7	Random Forest - Unbalanced	0.15474
8	Random Forest - Balanced	0.23570

References

1. O'Rourke, D. (2022, October 18). Brands losing a record \$29 for each new customer acquired. SimplicityDX. <https://www.simplicitydx.com/blogs/press-release-brands-losing-a-record-29-for-each-new-customer-acquired>
2. Sakar, C. & Kastro, Y. (2018). Online Shoppers Purchasing Intention Dataset [Dataset]. UCI Machine Learning Repository. <https://doi.org/10.24432/C5F88Q>.
3. Smith, B. (2024, May). How to improve each stage of your ecommerce sales funnel. wordstream.com/blog/ws/2020/07/01/ecommerce-sales-funnel